

Instituto Superior de Engenharia de Lisboa
BEng in Computer Science and Engineering
System Virtualization Techniques, Autumn/Winter 2024/2025
Third coursework assignment

For this assignment, we provide you with a minimal web application developed using Node.js, with data stored in Elasticsearch. The application supports running multiple instances simultaneously on different ports, anticipating the use of a reverse proxy with load balancing to distribute requests across these instances.

You will create a *systemd* service that receives control instructions on a Unix domain socket to:

- Launch one or more new instances of the Node.js application and add them to the reverse proxy.
 - Terminate one or more instances of the Node.js application and remove them from the reverse proxy.
 - Deactivate the application in the reverse proxy, stopping all instances while keeping the configuration.
 - (Re)launch all application instances and (re)activate the application in the reverse proxy.
-

Setup

Ensure that solutions for the **first** assignment are stored in the **cw1** directory of your repository, and that solutions for the **second** assignment are in **cw2**. If not, move the files now, then commit and push the changes.

Place *tvS-2425-1_cw3.tgz* in the **cw3** directory of your repository and execute: `tar xzvf tvS-2425-1_cw3.tgz`

After the extraction, delete *tvS-2425-1_cw3.tgz*, then commit and push the extracted files and directories.

Before starting this assignment, install additional packages in your Ubuntu 24.04 environment, by executing:

```
sudo apt update
sudo apt install npm nodejs nginx
```

Finally, download and install **elasticsearch** (choose the appropriate version for your CPU):

```
x86-64 ⇒ wget https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-8.15.3-amd64.deb
arm64 ⇒ wget https://artifacts.elastic.co/downloads/elasticsearch/elasticsearch-8.15.3-arm64.deb
sudo dpkg -i elasticsearch-8.15.3-*.deb
sudo nano /etc/elasticsearch/elasticsearch.yml
    xpack.security.enabled: false                                     <<< edit this configuration line
```

Preparation

- i. Confirm the proper operation of the provided web application:
 - a. Install the project's dependencies:

```
cd cw3/tvsapp/app
npm install
```
 - b. Launch only the web application, without the database:

```
NODE_PORT=43210 npm start
```
 - c. Use the browser to access <http://localhost:43210>
(you should see PORT: 43210 and *database unavailable*)
 - d. Start **elasticsearch**:

```
sudo systemctl start elasticsearch
```
 - e. Use the browser again to access <http://localhost:43210>
(you should see PORT: 43210 and a counter that increments with each new request made)
 - f. Terminate the web application (Ctrl-C), but keep **elasticsearch** running
- ii. Copy the web application from `cw3/tvsapp/app` to `/opt/isel/tvs/tvsapp/app` and set it up as a service. Then, start 4 instances in ports 43211, 43212, 43213, and 43214.
NOTE: see `cw3/tvsapp/etc/service/` and [Service Templates](#).

- iii. Add a configuration for a new site (tvapp) to /etc/nginx/sites-available listening on port 4444 and operating as a load balancer for the 4 local instances on ports 43211 to 43214. Activate the new configuration in /etc/nginx/sites-enabled along with the existing one (default). Use the browser to access <http://localhost:4444> and check the load balancer's operation (PORT changes on refresh).
NOTE: see cw3/tvsctl-srv/etc/nginx/sites-available/ and [nginx Configuration Control](#)
-

Exercises

- Write the following bash scripts to manage the configuration and operation of the proposed solution:
 - tvapp-reset.sh arguments: scale (default = 1), base (default = 35000)
Force an initial stopped configuration with scale instances in consecutive ports starting at base
 - tvapp-inc.sh arguments: delta (default = 1)
Add delta instances of Node.js running the web application, using more consecutive ports
 - tvapp-dec.sh arguments: delta (default = 1)
Remove delta instances of Node.js with the highest ports in use, leaving at least 1
 - tvapp-stop.sh arguments: -db (optional)
Deactivate the site from nginx and stop web app instances. If -db, also stop elasticsearch.
 - tvapp-start.sh
Start elasticsearch, all web app instances and (re)activate the site in nginx
 - tvapp-status.sh
Write a summary with the solution status, with one line per element (nginx, web apps, db)

You may use the following command to confirm that the load balancer is distributing the requests:

```
seq 32 | xargs -I{} curl -s http://localhost:4444/ | grep "PORT" |  
sed "s/<\/\?[a-z]\+>\/\/g" | sed "s/^\[:space:]*//" | sort | uniq -c
```

Tag this exercise on the GitHub repository with: **CW3-1** [Reference date: November 18th, 2024]

- Using the C language, build a *systemd* service (tvscld) to receive instructions on a Unix domain socket, to be placed in /run/isel/tvsctl/request, and a client program (tvsc1) with the following operations:
 - tvsc1 reset [scale [base]]
 - tvsc1 inc [delta].
 - tvsc1 dec [delta]
 - tvsc1 stop [-db]
 - tvsc1 start
 - tvsc1 status

The tvscld service is [socket activated](#) and runs with *root* privileges. Both the socket and the client (tvsc1) will only be available to *root* and members of group tvsgrp. to which user isel will belong.

The client program (tvsc1) is non-interactive and its first instruction must be close(0). For each execution of the client program, a request is sent to the service (tvscld) through the socket. The service *daemon* will then invoke the appropriate script (from exercise 1) to perform the requested action.

Write a script (cw3/install-all.sh) to copy all the files needed to run the solution from the repository folders to their final location (/opt/isel/tvs/...) and set up all permissions appropriately.

Tag this exercise on the GitHub repository with: **CW3-2**

ISEL, November 10th, 2024

Submission last date: November 28th, 2024